

THEME 2 · AI-POWERED DATA INTELLIGENCE PLATFORM

DataPilot AI

**Ask in plain English. Get validated,
explainable data intelligence.**

ASK

PLAN

COLLECT

VALIDATE

ANALYZE

EXPLAIN

THE PROBLEM

One data question shouldn't need a data team

"Compare EV sales in India from 2022 to 2025 and tell me which manufacturers grew fastest."

Today: 9 manual steps, needing SQL, Python and statistics

Find datasets > Download > Read columns > Clean > Dedupe > Fix missing > Join > Calculate > Chart

With DataPilot: one sentence

Type the question > Watch the workflow run > Read the dashboard + evidence

Not a dashboard

Fixed charts on one pre-loaded dataset. It can't answer a new question.


Not a chatbot

Answers from model memory. No real data, no validation, no source.

A data agent

Builds and runs a real workflow, then proves the answer.

Six stages from question to **auditable answer**




The workflow is dynamic

“Average price” needs collect → clean → average. “Rainfall vs crops” needs two sources, a join and correlation. “Sales spikes” needs time-series and anomaly detection. The orchestrator plans each one.

The workflow is transparent

The user watches each stage run live (✓ Understanding, ✓ Collecting, ● Validating, ○ Analyzing), so the system is never a black box.

An orchestrator coordinating specialist agents



Key design rule: the LLM plans and explains, but only calls a fixed toolset (`get_dataset`, `query_api`, `clean`, `validate`, `analyze`, `chart`). It never runs arbitrary code.

WORKED EXAMPLE

One sentence in, a full workflow out

1 · USER ASKS

“Compare rainfall and wheat production across Haryana districts from 2020 to 2025.”

2 · AI EXTRACTS INTENT

```
{  
  "location": "Haryana",  
  "period": [2020, 2025],  
  "variables": ["rainfall",  
               "wheat_production"],  
  "dimension": "district",  
  "task": "correlation"  
}
```

3 · AI BUILDS THE WORKFLOW

COLLECT Rainfall data + wheat production data

NORMALIZE District names and years

JOIN Merge on district + year

VALIDATE Rainfall 98.4% · Production 97.9% complete

ANALYZE Yearly trend, district comparison, correlation

VISUALIZE KPIs, line chart, district bars, scatter

EXPLAIN Insights linked to the source data

Validation first, then “Explain this result”

Validation Agent checks every dataset

Missing values

10,000 records · 125 missing · 98.75% complete

Duplicates

42 of 10,000 removed

Invalid values

EV sales = -500 is flagged

Types and units

“₹5 lakh” → number · MW / GW / kW normalized

Every number has a lineage trail

RESULT EV sales grew 34.8%

CALCULATION 2023 → 2025

DATASET EV Sales Dataset

CLEANING Deduped · names and units normalized

QUALITY 98.7% complete

SOURCE Public API / Cloudinary asset

Data quality is shown **before** conclusions, so bad data is never silently treated as reliable.

Cloudinary is our evidence layer



Uploaded datasets, screenshots, reference media and exported reports stay attached to the workflow that used them, giving the original asset a place in the provenance chain.

	Dashboard	AI chatbot	DataPilot AI
Handles new questions	No	Yes	Yes
Works on real data	Yes	Rarely	Yes
Validates data quality	No	No	Yes
Numbers computed deterministically	Yes	No	Yes
Shows where results came from	No	No	Yes, full lineage

TAKEAWAY

Don't make users learn data engineering.

User asks

Describes the question in plain English.

AI plans

Builds the workflow and coordinates the agents.

Code verifies

Deterministic validation and analysis produce the numbers.

Platform proves

Evidence and lineage sit behind every answer.

DataPilot AI. Thank you. Questions?